

XML2CSV-GENERIC-CONVERTER

VERSION

1.0.0

"EQUINOX"

AUTHOR

LAURENT POPIEUL

lochrann@rocketmail.com



Last update : Monday, September 21, 2015

Contents

1 - Background.....	4
2 - Audience.....	4
3 - XML2CSV-Generic-Converter.....	4
3.1 - XML2CSV: who cares?.....	4
3.2 - XML2CSV: how to?.....	5
3.3 - XML2CSV: what for?.....	6
3.3.1 - A few definitions, to settle things.....	6
3.3.2 - What does XML to CSV conversion exactly do?.....	6
3.3.3 - Optimization levels.....	8
3.3.3.1 - Raw level.....	8
3.3.3.2 - Standard optimization.....	8
3.3.3.3 - Extensive optimization – variant 1.....	10
3.3.3.4 - Extensive optimization – variant 2.....	11
3.3.3.5 - Extensive optimization – variant 3.....	11
3.3.4 - Filter files.....	13
3.3.4.1 - Positive filter file.....	13
3.3.4.2 - Negative filter file.....	14
3.3.5 - Structure analysis and XML processing.....	15
3.3.5.1 - Principle.....	15
3.3.5.2 - XML template file.....	15
3.4 - XML2CSV: what next?.....	16
4 - XML2CSVConsoleCommand.....	16
4.1 - Parameter list.....	17
4.2 - Return codes.....	20
5 - XML2CSVGenericGenerator.....	20
5.1 - XML2CSVGenericGenerator constructors.....	21
5.1.1 - File oriented constructors.....	21
5.1.2 - OutputStream oriented constructors.....	22
5.2 - XML2CSVGenericGenerator generation methods.....	22
5.3 - Other XML2CSVGenericGenerator methods.....	23
5.3.1 - Method setOutputStream.....	23
5.3.2 - Method setOutputFile.....	24
5.3.3 - Method setOutputDir.....	24
5.3.4 - Method setFieldSeparator.....	24
5.3.5 - Method setEncoding.....	24
5.3.6 - Method setOptimization.....	24
5.3.7 - Method setCutoff.....	24
5.3.8 - Method setWarding.....	24

5.3.9 - Method setUnleashing.....	25
5.4 - XML2CSVGenericGenerator exceptions.....	25
5.5 - LoggingFacade abstract class.....	25
6 - Command line examples.....	26
6.1 - Plain conversion of one XML file.....	26
6.2 - Extensive #1 conversion of several XML files with attributes.....	26
6.3 - Extensive #3 conversion of one XML file with log.....	27
6.4 - Extensive #2 conversion of one XML file with many options.....	27
6.5 - Raw conversion of several XML files in a blend output file.....	27
7 - Real life XML examples.....	28
7.1 - All's Well That Ends Well.....	28
7.2 - 100 MB experiment.....	28
7.3 - The loop is closed now.....	29

1 - Background

Date	Author	Version	Description
2014-03-21	Laurent Popieul	1.0.0 ("Equinox")	From scratch.
2014-04-21	<i>ditto</i>	<i>ditto</i>	Adjustments.
2014-05-21	<i>ditto</i>	<i>ditto</i>	Most complete optimization mode added.
2014-06-21	<i>ditto</i>	<i>ditto</i>	Finalization.
2014-12-21	<i>ditto</i>	<i>ditto</i>	Proofreading.

Written with Open Office Writer 4.1.1 using Cédric Billard's documentation model available at <http://templates.openoffice.org/fr/template/documentation>.

2 - Audience

XML2CSVGenericConverter user documentation - Copyright 2014 Laurent Popieul.

This document describes what the program does, how to execute it, and its runtime parameters.

The targeted audience consists of programmers and advanced end users.

XML2CSVGenericConverter and this accompanying documentation come with ABSOLUTELY NO WARRANTY. This is free software, and you are welcome to redistribute it under certain conditions.

3 - XML2CSV-Generic-Converter

3.1 - XML2CSV: *who cares?*

XML has become a standard for data exchange between heterogeneous systems but it's quite a verbose language and some files grow beyond reason.

Fortunately XML files can be zipped easily for transport and computer experts know how to build efficient programs which can deal with XML monsters.

From time to time however, end-users need to pick up data from XML files by hand and check them.

For instance:

- because data were expected in the last XML files but can't be found anywhere in the IS, and though the IT keeps telling us everything is OK and no error occurred, someone has to check things out explicitly;
- because the program displayed funky figures and the calculation has to be double-checked by hand from the XML source.

Small XML files are no problem at all, and average size XML files can be imported by a spreadsheet and easily reworked from it.

Trouble comes with big XML files and it's not uncommon to see products fail, freeze, crash or collapse when confronted to XML of a respectable size.

There are two approaches in such situations:

1. to have the file opened by a smart editor (such as Notepad++) and then filter/clean the data in order to discard everything but the meaningful tags (for instance with carefully devised *regular expressions*);
2. or to rely on a utility program which will extract data from this nasty hotchpotch.

I wrote this XML2CSV converter because I was not too pleased with option 1. 's error prone manipulations, and not exactly satisfied with what I had up my sleeve so far to handle option 2.

I wanted something generic which would basically convert all XML data to CSV for later opening with a spreadsheet without bugging me with things like enclosing blocs or repeated elements.

XML is simple, CSV is simple, so why would XML \Rightarrow CSV have to be complicated?

I wanted some extra options too, making it possible either to select some XML tags and have the program extract them only, or to select some XML tags and have this time the program extract everything but those ones.

I must admit that the resulting program was bigger in the end than what I expected in the first place. Indeed, placing cells at the right place can be a tricky job sometimes, especially when a new XML tag might pop up anytime because you have no XML schema telling you what is possible and what is not.

Anyway, I was personally satisfied enough with the result – at least for my own needs – and decided to share it with other fellow XML consumers who might look for something alike.

3.2 - XML2CSV: *how to?*

The generic converter is packaged as an executable Java Jar file.

The core Java class is named `XML2CSVGenericGenerator` and holds a `generate` method in charge with both XML \Rightarrow CSV conversion and CSV output file generation.

A wrapping class named `XML2CSVConsoleCommand` is responsible for interactions with the end user and, just like its name suggests it, in charge of console command execution.

This wrapping class is declared in the Jar file as the executable class of the package which means that when the Jar is executed `XML2CSVConsoleCommand` is run automatically.

In other words, showing you how to run XML2CSV conversion goes back to show you how to execute a Java Jar with or without parameters. In our case, it amounts to:

- copy `XML2CSVGenericConverter_V1.0.0.jar` to a local directory on your machine;
- open a console window with Java runtime resources available – 1.6 or higher ⁽¹⁾;
- go to the directory where you copied `XML2CSVGenericConverter_V1.0.0.jar`;
- type `java -jar XML2CSVGenericConverter_V1.0.0.jar` and press “enter” to execute the Jar without parameter (you will get a WARNING/ERROR message telling you that parameter `-i` is missing ⁽²⁾, and another one inviting you to rerun the program with option `-h` in order to display the on-line help);
- type `java -jar XML2CSVGenericConverter_V1.0.0.jar -h` and press “enter” to execute the Jar with parameter `-h`, in order to display the on-line help;
- type `java -jar XML2CSVGenericConverter_V1.0.0.jar` followed a list of parameters (complete list provided at § 4.1 -) and press “enter” to execute the Jar with the options you chose.

¹ Type: `java -version` and press “Enter”. If the system is well configured you should see a message displaying the current Java version on your system. If you see an error message then you should make Java available on your system first before you try and perform XML2CSV conversion again.

² The only mandatory parameter.

If you are a more advanced Java user then you might prefer to execute the Jar's main class yourself with the following command:

```
java utils.xml.xml2csv.XML2CSVConsoleCommand {parameters}
```

where {parameters} represents your runtime parameters in sequence (see § 4.1 -).

Don't forget to add `XML2CSVGenericConverter_V1.0.0.jar` to your local Java `CLASSPATH` before.

3.3 - XML2CSV: *what for?*

3.3.1 - A few definitions, to settle things

In an XML structure a tag might contain other tags or data, and have attributes.

In order to settle things I will call in the rest of this document a tag containing plain data a *leaf* tag, and a tag containing other tags a *composite* tag, *parent* of the *children* tags enclosed in it.

An *ancestor* is the *parent* of a *parent* tag, or the *parent* of that *parent*, and so on. The *root* tag, which is the top level tag which contains all the other tags, is the *parent* or an *ancestor* of all tags.

A *composite* tag can contain *composite* tags, *leaf* tags, or a mix.

Composite and *leaf* tags can be *repeated* in a valid XML file, or remain *single*. A *repeated* tag will appear at least twice within the same *parent* tag.

A tag which does not appear systematically at its position in a *repeated composite* tag is named an *optional* tag.

In the sample XML file below `<Row>` is a *composite repeated* tag as well as the `<Data>` tag within a `<Row>`.



sample.xml

The `<Other>` tag inside a `<Row>` is a *composite* tag too but it is not *repeated* in each `<Row>`.

The `<Data0>`, `<Amount>` and `<Data1>` tags are *single leaf* tags inside the *composite* `<Data>` tag and the three of them are *optional* tags because they do not always exist in each `<Data>` tag ⁽³⁾.

The `<Date>` tag inside a `<Row>` is a *repeated optional leaf* tag.

Two tags are *independent* when:

- neither of them is the *parent* or *ancestor* of the other, and,
- they don't have any common *parent* or *ancestor* bearing a *leaf* tag or attribute, *root* tag excluded.

In the sample XML file the `<Head>` tag and the `<Row>` tags are *independent*, `<Row>` tags are *independent*, but none of the sub tags inside a `<Data>` tag are *independent* because the `<Data>` tag holds at least one *leaf* tag.

3.3.2 - What does XML to CSV conversion exactly do?

The XML ⇒ CSV converter generates one column for each XML *leaf* tag.

The content of each *leaf* tag is put in a row in the corresponding column, one under the other if it is a *repeated leaf* tag.

3 In fact, all leaf tags in the sample XML file are optional.

Data which are *directly* or *indirectly related* might appear on the same line for a more compact representation, or left in separate rows.

Here is my definition of *related* elements:

- if you consider one initial XML *single leaf* tag and if you follow its *parents* back to a certain enclosing level in the XML file then each *single leaf* tag which is also a *child* of one of those chained *parents* and *ancestors* is *indirectly related* to the initial XML *leaf* tag, provided that all the *parents* of the chain are *single* too (apart from the last enclosing *parent* which might be *repeated* or not);
- if you consider one initial *single leaf* XML tag, all the other XML *single leaf* tags which have the same *parent* tag are *directly related* to that initial tag (no matter if the *parent* is *repeated* or not).

Relationship and *dependency* (as defined in § 3.3.1 -) are distinct concepts but they are consistent with each other: when two *leaf* tags are *related* all the data of their shared enclosing *parent* or *ancestor* are *dependent* (but the reverse is not true).

If you look at the `<Data0>` tag inside the `<Data>` block of the 1st `<Row>` in the previous file `sample.xml` you will see that it is closely *related* to the `<Amount>` and the `<AddAmount>` elements of that 1st `<Row>`, so `true`, `123` and `12` might appear on the same CSV output line. All the tags in the `<Data>` block are *dependent*, but not all of them are *related*.

Grouping *directly* and *indirectly related* tags on the same line can be interpreted as linking them to the same *adoptive parent*.

Once *related* tags have been grouped on the same line they might be merged back into lines associated with their *half-brothers* (that is, *repeated leaf* tags which have the same *parent* as the *adoptive parent* of these *related* tags) or even more remote *half-distant-nephews* (*repeated leaf* tags which have the *adoptive parent* as *ancestor*).

Directly related fields may be merged back into lines associated with their *distant-nephews* too (that is: *single leaf* tags which have the *parent* of these *related* tags as *ancestor*).

You might also safely consider merging *related* tags back into lines associated with *single* or *repeated leaf* tags which have the *adoptive parent* as *ancestor* provided that at least one of these chained *parents* is *repeated*.

In the previous file `sample.xml`, `true`, `123` and `12` might be duplicated back on the lines displaying `ee1` and `ee2`, and on the lines displaying `SD1a` and `SD1b` as well.

The CSV output file could look like the `sampleOutput.csv` file below, but it could be slightly different too, depending on your chosen packing/duplicating strategy.



sampleOutput.csv

However, whatever the strategy, *repeated* tags are never *related* with one another and, as such, can never be mixed and must always appear on separate lines:

- one line under the other for the same *repeated* tag;
- in different line blocks for distinct *repeated* tags (no matter if they have the same *parent* or not).

If you think of XML *leaf* tags in their corresponding CSV lines then the ultimate goal when it comes down to XML ⇒ CSV conversion is to have for each of them all its kin present on the same line, and the fewest number of lines as well. When the two conditions are met "utmost compactness" is reached and it becomes possible to process separately one line at a time from a resulting file holding no more CSV lines than needed... A bit of lightness in a world of heavies.

XML2CSVGenericConverter offers five different packing/merging strategies, one for each optimization flavor available at runtime. They are described in the next paragraphs, from the simplest to the most sophisticated.

3.3.3 - Optimization levels

3.3.3.1 - Raw level

Raw optimization level stands for no optimization at all, and no data packing at all.

In this mode the content of each *leaf* element is put alone in a new CSV line at its expected column index, and the output CSV file will hold as many lines as there were XML *leaf* elements.

Because no data buffering of any kind is needed for such an elementary behavior, the data & size of the XML input file(s) have absolutely no impact over performance even in nasty scenarios such as the one implying a heavily *repeated leaf* tag (say: `<paininthe>BACK</paininthe>` repeated thousands of time in a row and you'll get it).

The file below is the raw XML \Rightarrow CSV conversion of the previous `sample.xml` file generated by XML2CSVGenericConverter:



sampleRawOutput.
csv

3.3.3.2 - Standard optimization

Standard optimization consists in packing *directly related* fields on the same CSV line at each level of the tag hierarchy.

This mode produces fairly compact CSV output files in as much as:

- each XML element content appears once only in the output file(s);
- packing *directly related* fields reduces the number of lines.

Data layout in the CSV file(s) is also very much like what it is in the original XML counterpart.

For all these reasons, I chose it as the default optimization triggered when no particular optimization directive (`-r` or `-x`) is provided on the command line.

The file below is the standard XML \Rightarrow CSV conversion of the previous `sample.xml` file generated by XML2CSVGenericConverter:



sampleStandardOut
put.csv

Field packing is performed in memory and requires that a consistent amount of data be buffered, that is, to wait until the last XML closing tag of the last *parent* of *leaf* tags is read.

Each time a buffer reaches consistency it is packed, then flushed to the CSV output, and buffering starts fresh anew.

This approach helps to keep the buffer fairly small in as much as *independent* XML blocks (see § 3.3.1 -) are processed by different data buffers (because when the last tag of an *independent* XML block is read the buffer always reaches consistency and is flushed).

In the previous `sample.xml` file the `<Head>` block and each `<Row>` blocks are *independent* and no more than one `<Row>` block is ever loaded in memory at a time by XML2CSVGenericConverter.

Unfortunately the reverse is also true: *dependent* XML blocks (that is: blocks which have at least one *parent* or

ancestor in common bearing a *leaf* tag or attribute) will be loaded by the same data buffer waiting to reach consistency, and if it happens that one of them is heavily repeated then the risk of `OutOfMemoryError` ⁽⁴⁾ cannot be ruled out.

Imagine that you change the previous `sample.xml` file in order to enclose all the `<Row>` blocks in one `<Rows>` tag, and that you add a `<Total>` *child leaf* tag to `<Rows>` and/or add a `total` attribute to `<Rows>` (provided that you decided to have attributes extracted by adding `-a` on the command line). By doing this you would create *dependent* `<Row>` blocks which would be processed by the same data buffer. Needless to say that if the XML file contained thousands of `<Row>` blocks instead of three it might lead to a serious runtime memory problem.

The `sample2.xml` file below is an illustration of the previous statement:



sample2.xml

If you execute `XML2CSVGenericConverter` against this file in debug mode (with option `-d`) you will see that the program loads the whole contents of the three `<Row>` blocks in the same buffer to reach consistency, something which doesn't occur with the original `sample.xml`.

That's where negative *filter files* come into the act (see §. 3.3.4.2 -). They make it possible to exclude tags from the program scope and happen to be very helpful to reduce data *dependency* and actually convert big XML files after a frustrating initial "Java heap space" issue (`OutOfMemoryError`).

Sometimes, excluding one unique tag between two loosely *dependent* XML blocks (provided that this tag is not interesting and might be cast aside) reduces dramatically the overall data buffer size and makes the difference between failure and success.

For example, `XML2CSVGenericConverter` execution against the previous `sample2.xml` file with a *negative filter file* excluding `Root.Rows.Total` from the conversion (and `Root.Rows@total` if you have attributes extracted as well) makes it possible to parse each `<Row>` in a different data buffer again, and, just like with `sample.xml`, in a much more memory friendly way.

The bad news is that data *dependency* level varies between XML flavors and that beyond a certain *dependency* threshold and a certain XML size no program, not even my `XML2CSVGenericConverter`, will achieve the task.

The good news is that `XML2CSVGenericConverter` offers you a way to reclaim an active control over the XML *dependency* level without actually changing a byte in the original files.

I chose to ship `XML2CSVGenericConverter` with a built-in *root* tag "exemption" behavior which is active by default (but which might be deactivated at your will with option `-u`). The idea behind the scene is that for the vast majority of real life XML files the *root* tag holds special attributes which artificially magnify data *dependency* (name spaces alias declarations for instance) and that would have meant to systematize *negative filter file* usage to comply with that, something I didn't like. When you don't like reality, change it! So I coded a inner "root gambit" which automatically excludes the 1st level of the tag hierarchy from packing and, as such, nullifies *root*'s artificial *dependency* burden at the cost of inducing a "root edge effect" which prevents efficient packing optimization at the very 1st level of the tag hierarchy.

From a practical point of view if you deal with small XML files only then *root* tag "exemption" deactivation (by adding `-u` on the command line) will be affordable. SAX parsing won't probably be any more memory-friendly than well known DOM-like approaches (because the whole XML file will most certainly have to be buffered for optimization consistency) but you won't get any "root edge effect".

But remember: `XML2CSVGenericConverter` was devised in the first place to handle big files so "root gambit"

4 A "Java heap space" issue.

will be the rule and option `-u` the exception, left at your clean-sighted responsibility.

To be exhaustive on this question advanced users should be aware that two more atypical XML flavors would also potentially lead to memory overload, namely something like:

- an XML file bearing a flock (say: thousands) of different *leaf* tags inside the same *parent* tag;
- an XML file bearing a heavily *repeated leaf* tag (again, thousands or more).

Here again, the buffer would grow accordingly to fit the data before the program actually gets a chance to flush it (waiting to reach consistency) at the risk of an `OutOfMemoryError` (in the 1st case: caused by a outrageous amount of buffer columns; in the 2^d case: caused by an outrageous amount of buffer lines).

Fortunately this time, it is generally considered very very bad practice (and I just can't imagine which dark purpose it could serve to compose such an alien XML).

However, should anybody come face to face with hostile Xenomorphic Markup Language, the previous raw optimization level would provide a quick fix and still produce some operable output (see § 3.3.3.1 -).

3.3.3.3 - Extensive optimization – variant 1

In its first variant, extensive optimization consists in:

- packing *directly related* fields on the same CSV line, at each level of the tag hierarchy, just like the previous optimization does;
- at each level of the tag hierarchy, replicating the previous *related* fields back into lines associated with *half-brothers* (CSV lines which are associated with the content of a *repeated leaf* tag having the same *parent*) or *half-distant-nephews* (CSV lines which are associated with the content of a *repeated leaf* tag having the same *ancestor*);
- replicating the *directly related* fields back into lines associated with their *distant-nephews* (CSV lines associated with the content of a *single leaf* tags having the same *ancestor*).

In other words extensive optimization - variant 1 - is a standard CSV optimization of *single leaf* tags plus copy back into CSV lines representing *repeated leaf* tags of the same XML level or deeper level in the same XML block, plus copy back into deeper *single leaf* linked lines in the same XML block.

The file below is the extensive #1 XML ⇒ CSV conversion of the previous `sample.xml` file generated by XML2CSVGenericConverter:



sampleExtensiveV1
Output.csv

This extensive optimization uses the same buffering strategy as the one used for standard optimization (see § 3.3.3.2 -) with the same features/limitations concerning *dependent* data.

Extensive optimization #1 may not produce the best XML ⇒ CSV conversion results (that is: the less CSV lines possible) because its standard packing strategy handles *directly related* fields only and casts aside *indirectly related* fields (see § 3.3.2 -).

Another extensive optimization, based this time on both *directly* and *indirectly related* fields, is described in the next paragraph.

3.3.3.4 - Extensive optimization – variant 2

In its second variant, extensive optimization consists in:

- packing *directly* or *indirectly related* fields on the same CSV line, at each level of the tag hierarchy, associated with the same *adoptive parent*;
- at each level of the tag hierarchy, replicating the previous *related* fields back into lines associated with *half-brothers* (CSV lines which are associated with the content of a *repeated leaf* tag having the same *parent*);
- at each level of the tag hierarchy, replicating those *related* fields back into lines associated with *single* or *repeated leaf* tags which have the *adoptive parent* as *ancestor* provided that at least one of these chained *parents* is *repeated*.

The file below is the extensive #2 XML \Rightarrow CSV conversion of the previous `sample.xml` file generated by XML2CSVGenericConverter:



sampleExtensiveV2
Output.csv

It looks pretty much like what extensive #1 XML \Rightarrow CSV conversion does, which is not very surprising (and maybe quite reassuring too).

However, extensive #2 XML \Rightarrow CSV conversion might produce better results (that is: less CSV lines) than extensive #1 XML \Rightarrow CSV conversion with XML containing a lot of *single composite* tags and a lot of *single leaf* tags.

This extensive optimization uses the same buffering strategy as the one used for standard optimization (see § 3.3.3.2 -) with the same features/limitations concerning *dependent* data.

Paradoxically, the efficiency of extensive #2 optimization increases with the data *dependency* degree and is maximum when the data issued from an XML block bears a top *leaf* element and/or a top attribute, for it makes it possible to “fold” at best the whole data block. The paradox is only apparent because the more data are *related*, the more data packing will be efficient, but the more the corresponding XML blocks will be *dependent* and thus increase the buffer size.

In other words, high quality XML \Rightarrow CSV optimization is a matter of compromise between efficiency and memory preservation. The last extensive optimization, which is described in the next paragraph, takes fully advantage of this statement.

3.3.3.5 - Extensive optimization – variant 3

Extensive #3 XML \Rightarrow CSV conversion uses the very same optimization algorithm as extensive #2 XML \Rightarrow CSV conversion.

The gimmick here consists in adding virtual hidden attributes to carefully selected XML tags beforehand which will maximize data “folding” without compromising memory usage, and then to proceed with extensive #2 XML \Rightarrow CSV conversion.

Those virtual attributes are nothing more than reaction catalysts: not important for themselves but yet crucial because they dramatically increase reaction performance by their mere existence.

And this is exactly how it goes.

Technically, the program builds an inner list of all XML tags which:

- are either *repeated* “enough” (conventionally: at least 10 times through the entire XML file),
- or are not *repeated* “enough”, but in addition, do not contain any inner repeated “enough” tag (*child* or

descendant).

Those criteria are important because they ensure that data *dependency* doesn't explode and cause memory wreckage. Number 10 doesn't have any mysterious & magical property as far as I know, but it works. 5 or less is not enough and 15 or more doesn't change anything so I guess any number above 5 is OK... Maybe a mathematician could help us on the topic but in the meantime I go on with a hard coded 10 ⁽⁵⁾.

Tags which have already one attribute are excluded from the list, and the *root* tag is also excluded (for performance issues). The remnant tags receive a virtual hidden attribute value during data buffering and classical extensive #2 XML \Rightarrow CSV conversion is performed as usual.

A virtual hidden attribute is never sent to the output. It is silently discarded after optimization phase when its "folding" job is over just like a regular catalyst.

The three files below illustrate this principle: `policies.XV2.csv` and `policies.XV3.csv` are respectively the extensive #2 XML \Rightarrow CSV and the extensive #3 XML \Rightarrow CSV conversion results of `policies.xml`.



policies.xml



policies.XV2.csv



policies.XV3.csv

The effectiveness of aggregation catalysis is obvious.

Because the *root* element is excluded from the aggregation process (remember the "*root* gambit" explained at § 3.3.3.2 -) it prevents the data extracted from the `<Sender>`, `<FileData>`, `<Validity>` and `<EmptyIndic>` blocks from being aggregated with the data from the `<Policy>` blocks, but "folding" is maximum in all the rest of the hierarchy.

However, because `policies.xml` is a rather small XML file *root*-inclusive optimization is not a nonsense...

So, for those into "extreme packing" who might have been a bit frustrated by the previous `xv3` CSV file the `policies.unleashed.XV3.csv` file below is the *unleashed* extensive #3 XML \Rightarrow CSV conversion result of `policies.xml` (that is: performed with an additionnal `-u` option in order to include the *root* tag in the packing plan).



policies.unleashed.
XV3.csv

By the way: in `policies.xml` the *root* element `<Policies>` contains one unique *leaf* element `<EmptyIndic>` which seems harmless but which is in fact solely responsible for a maximization of the global data *dependency* level of that XML file. The previous *unleashed* conversion works well because `policies.xml`

5 I have so far resisted the idea to add an uncanny `-z` (for *shazam!*) command line option parameterizing the figure but the temptation is, hum, still pretty much tantalizing.

contains three `<Row>` blocks only, but had it contained three thousands `<Row>`s it would have been another story.

However, because the *root* element is excluded from “folding” with the default settings there is no need to exclude `<EmptyIndic>` from conversion with a dedicated *negative filter* file to have `<Policy>` blocks become *independent* again, and perform memory-friendly conversion no matter how many policies are listed.

3.3.4 - Filter files

3.3.4.1 - Positive filter file

When left to itself, `XML2CSVGenericConverter` assumes that all XML *leaf* tags are interesting and should appear as columns in the CSV output, and behaves accordingly.

This said, the end user may select an explicit subset of XML *leaf* tags, gather them in an appropriate file listing the corresponding XML element Xpaths and rerun `XML2CSVGenericConverter` with option `-p` to have the program extract those fields only (see § 4.1 -).

This file, which I named a *positive filter file* for obvious reasons, is expected UTF-8 encoded by the program without Byte Order Mark (BOM).

It might contain comments, that is, lines beginning with two consecutive dashes (`--`), and might contain blank lines as well.

Xpaths have to be listed one under the other (one per line), and the program will smoothly keep the same order for the generated columns in the output file(s).

Duplicate Xpaths are sorted out automatically.

Junk Xpaths which do not exist in the XML input (⁶) are discarded, and an appropriate warning message is displayed (but the program will stop if all the Xpaths listed in the *filter file* happen to be incorrect).

Xpaths of *composite* tags might be mixed with Xpaths of *leaf* tags, and the program will automatically replace a *composite* Xpath with all the *leaf* Xpaths it contains, eliminating duplicates.

For instance the *filter file* below, consistent with our initial `sample.xml` XML file, contains three *leaf* Xpaths (`Root.Row.Data.Data0`, `Root.Row.Data.Amount` and `Root.Row.Date`) plus one *composite* `Root.Row.Other` Xpath that the program replaces automatically with the *leaf* Xpaths it contains (namely `Root.Row.Other.Data.Data1` and `Root.Row.Other.Data.Data2`).



sampleFilterFile.txt

Applying this *filter file* in a *positive filtering* context on `sample.xml` would extract the `Data.Data0`, `Data.Amount`, `Date`, `Other.Data.Data1` and `Other.Data.Data2` tag values only and put them in a CSV output file.

The file hereafter is the extensive #1 XML ⇒ CSV *positive* output generated by `XML2CSVGenericConverter` against `sample.xml` when plugged onto the previous `sampleFilterFile.txt` with option `-p`:

⁶ I'd rather say: which do not exist in the in-memory structure representation of the XML input that the program builds on the fly. This remark is not as mundane as it seems, because, as you will read in § 3.3.5 - , the program just ignores data which do not comply with its inner structure representation, solely build out of one of the XML input files.



samplePositiveOutput.csv

Filter files allow attribute filtering (provided that you decided to have attributes extracted by adding `-a` on the command line) which means that Xpaths listed in a regular *positive filter file* - or a *negative* one, see below - might represent XML elements or attributes.

You will recognize an element's attribute from the column header name in a CSV output file. Indeed, a `@` separation character replaces the dot between the attribute name and its enclosing element Xpath (for instance: `Root.Row.Data.Amount@Currency`, if `Currency` is an attribute of the `Amount` tag).

This Xpath representation is also what the program expects in *filter files* to feature an attribute: you might put `Root.Row.Data.Amount` and `Root.Row.Data.Amount@Currency` in sequence in a *filter file* and run the program with option `-p` to extract both the corresponding amount and its currency.

I chose to display attributes by default in columns placed after the one dedicated to their enclosing element, just like spreadsheets do.

This said, because you have control over the column order in *filter files* you might choose to reverse this default order at your will (with the previous example: it you put `Root.Row.Data.Amount@Currency` before `Root.Row.Data.Amount` in your *filter file*) or even separate an element from its attributes.

I didn't want to enclose attribute support in version 1.0.0 but finally gave it better thoughts and shipped the product with minimal support for *leaf* elements first, and now with full attribute support (⁷).

3.3.4.2 - Negative filter file

Negative filtering is just the opposite of *positive filtering* (see § 3.3.4.1 -).

A *negative filter file* has the same structure as a *positive filter file* (UTF-8 encoded without BOM, *composite* or *leaf* Xpaths in columns, and so on) but when fed with a *negative filter* with command line option `-n` (instead of `-p`) the program extracts all XML fields but those listed in the *negative filter file*.

The file below is the extensive #1 XML \Rightarrow CSV *negative* output generated by XML2CSVGenericConverter against `sample.xml` when plugged onto the previous `sampleFilterFile.txt` with option `-n`:



sampleNegativeOutput.csv

I must admit that I added this option in the first place just for plain symmetry considerations with *positive filtering* without really thinking of any clever use of it.

The influence of data *dependency* on data buffering (see § 3.3.3.2 -) sheds now new light on *negative filtering*. I see it now as an important *dependency* downsizing facility which really comes in handy in order to successfully convert big XML files nesting heavily *dependent* data.

⁷ The effect of spring sunlight on a tired soul is just amazing.

3.3.5 - Structure analysis and XML processing

3.3.5.1 - Principle

`XML2CSVGenericConverter` can process one or several files in a row. Indeed, option `-i` might indifferently point to a single XML file, or to a directory ⁽⁸⁾ from which the program will process all XML files ⁽⁹⁾ in sequence.

XML input files processed consecutively are expected alike, that is, sharing the same general XML structure.

This is important because `XML2CSVGenericConverter` performs an in-memory analysis of that structure before the data of the first XML file is processed, and confidently relies later on this structure representation to fetch data.

In details, when structure analysis phase is triggered the program picks up one of the XML input files at random and parses it, not in order to retrieve its data but to look at the tag hierarchy imbrication (figuring *composite* and *leaf* tags), the tag cardinalities (figuring *single* or *repeated* tags), and the data types as well (something which has no real purpose so far but who knows what could come up with version 1.1.0?..).

3.3.5.2 - XML template file

I name the XML file which serves as model during the structure analysis phase the XML *template file*.

The random default *template file* picking lottery described in the previous paragraph might be bypassed anytime using the `-t` command line option, which makes it possible to provide the name (without its path) of the particular XML input file which should be used as *template* among those listed with option `-i`.

I added this option to let advanced users who have comparable files to process and who know that one particular XML input file is the most complete (I mean: with all its optional tags appearing at least once) help the program pick a winning lottery *template* ticket.

It might be useful because the lenient approach I've implemented, which ignores any tag that doesn't belong to the in-memory XML structure representation, has a direct & definitive consequence over optional tags which pop up unaware out of the scope of the *template file*: they are just plainly ignored.

To illustrate this, imagine that you have two XML input files `myFirst.xml` and `mySecond.xml` in the same input directory, and one *optional single leaf* tag – say `Root.Row.Unexpected.Data.Tag` – appears in `mySecond.xml` but not in `myFirst.xml`. Imagine now that you run `XML2CSVGenericConverter`, and that the program chooses `myFirst.xml` as *template file*. As a result, `Root.Row.Unexpected.Data.Tag` will be completely ignored by the program which won't even acknowledge it for an acceptable *positive* or *negative filter* Xpath.

Of course, running `XML2CSVGenericConverter` N individual times (one for each XML input file) is a good workaround because when a single XML input file is processed *template* and actual data are altogether the same. But it might become somewhat tedious when you have dozens of XML input files to process.

So, because I didn't want to systematically parse all XML input files just to make sure I would not miss any overdue hidden *optional* tag (a sheer waste of time) I ended up with another option, `-t`, and this paragraph.

8 Absolute or Jar-relative paths are welcomed.

9 That is, files with a `.xml` extension.

3.4 - XML2CSV: *what next?*

Raw and standard optimizations are the very least any decent XML ⇒ CSV conversion program might do, and just what most users will ask from it.

The extensive #1 and #2 optimization flavors (respective codes: `xv1` and `xv2`, see § 4.1 -) are interesting stepping stones which helped me devise the last and most complete extensive #3 optimization mode (code `xv3`).

However, you should keep in mind that the more elaborate the optimization is, the slower the conversion will be.

Extensive #3 optimization is nothing more than regular extensive #2 optimization of data previously “powdered” with additional virtual attributes catalyzing maximum data packing, but at the cost of roughly doubling the overall process time. Extensive #2 optimization or extensive #1 optimization might be your best choice when you don't need maximum compactness but look after speed. As for standard optimization it would be a serious mistake to scorn it, for it performs meteoric XML ⇒ CSV conversion.

From now on, I don't really expect to create an additional & hypothetical #4 variant because the latest testing I've performed to compare `xv3` against spreadsheets' XML uploads reassured me about my algorithm's reliability (see § 7.3 -).

Even though it is quite ugly in my opinion, the XML standards make it possible for an element to be both *leaf* and *composite*, that is, to loosely contain both plain data and other elements just like in the excerpt below:

```
<text>Mr. <name>Smith</name>'s order (number <orderid>123</orderid>) has been shipped.</text>
```

As a rule of thumb loose data should be avoided in a structured language such as XML.

As for XML ⇒ CSV conversion there is no definitive way to render loose data properly in a CSV file.

You should be aware that `XML2CSVGenericConverter` will treat a “mixture” element as a regular *composite* element holding an extra terminal *leaf* part, and will ignore all the rest. With the example above the program would extract “Smith” (element `name`) and “123” (element `orderid`) from `text` seen as a *composite* element, plus “) has been shipped.” seen as an extra terminal *leaf* part of `text`. Another example is provided at § 7.1 - .

Regarding all the other options added I implemented them one by one while I was writing this documentation until I had no more ideas of features to add, no more paragraph to remove from § 3.4 - , and no time left either.

I'm lying, you know?

I was just running out of `GetOpt` command line letters – that's all. Ho ho ho.

Anyway: some of those options might happen to be helpful.

Sometime.

4 - XML2CSVConsoleCommand



`XML2CSVConsoleCommand` is the main class of the executable Jar package and makes it possible for end users to interact with the underneath core & versatile `XML2CSVGenericGenerator` class (see § 3.2 -).


The paragraph below lists all the command line parameters which are supported by `XML2CSVConsoleCommand` making it possible to tailor execution to your specific needs.



Invoking `XML2CSVConsoleCommand` with an unknown parameter, or a supported parameter but in an


unsupported way will lead to hasty program termination to avoid unexpected behavior (\Rightarrow *graceful failure*).

4.1 - Parameter list

Parameter	Description
-h	<i>Help.</i> This option displays the on-line help and then aborts the program immediately no matter what the other parameters might be. Optional.
-m	<i>Mute.</i> When this option is set nothing is displayed at runtime and the program's return code only might indicate success/failure. See § 4.2 - . Optional.  Option -m should be placed at the head of the parameter list in order to be fully effective. <i>Option -m overrides options -v and -d.</i>
-v	<i>Verbose.</i> When this option is set extra progression messages are displayed. Optional.
-d{degree}	<i>Debug.</i> When this option is set debug messages are displayed. Optional. Option -d can be immediately followed (that is: without any gap) by an explicit degree value (1, 2, or 3) or provided alone (which transparently defaults to 1); the higher the debug degree is, the more debug messages will be displayed at runtime.  Warning: debug messages are extremely time & space consuming and should not be activated when big XML files are processed (it could slow down computation by at least two orders of magnitude and generate Godzilla sized log files).
-a	<i>Attributes.</i> Extracts elements' attributes as well. Optional.
-r	<i>Raw.</i> Optimization deactivation. Generates CSV output file(s) in which each extracted XML tag is placed alone in its own CSV line (see § 3.3.3.1 -). Optional. <i>Options -r and -x are exclusive.</i>
-x{variant}	<i>Extensive.</i> Elaborate optimization. Generates CSV output file(s) where related data are packed on the same line (see § 3.3.3 -) and eventually duplicated on repeated elements' lines. Optional. Option -x can be immediately followed (that is: without any gap) by an explicit optimization variant name (xv1, xv2 or xv3) or provided alone (which transparently defaults to xv3). xv1 is the command line name of extensive optimization – variant 1 (see § 3.3.3.3 -). xv2 is the command line name of extensive optimization – variant 2 (see § 3.3.3.4 -).

Parameter	Description
	<p>xv3 is the command line name of extensive optimization – variant 3 (see § 3.3.3.5 -).</p> <p><i>Options -x and -r are exclusive.</i></p>
-i {input dir/file}	<p><i>Input.</i> Path to an XML file or path to a directory containing XML files to convert to CSV. The path might be relative or absolute. A relative path is calculated against the local Jar directory.</p> <p> A path containing blank characters has to be enclosed by double quotes (example: -i ". /xml files/myFile.xml").</p>
-t {template filename}	<p><i>Template.</i> This option makes it possible to choose which of the input XML files (listed by the previous option -i) will serve for structure analysis (see § 3.3.5 -). A correct <i>template filename</i> consists of the file name only without its path. Optional.</p> <p><i>The program picks up at random one of the input XML files when left to itself.</i></p>
-o {output dir}	<p><i>Output.</i> Path to an existing directory where CSV output file(s) will be generated. Optional.</p> <p><i>The program uses the input directory when left to itself.</i></p>
-b{blend filename}	<p><i>Blend.</i> When this option is set the program generates a single CSV output file instead of generating as many CSV output files as XML input files (keeping names, changing extensions to CSV). Optional.</p> <p>Option -b can either be provided alone or immediately followed (that is: without any gap) by a <i>blend filename</i>. In the former case the program will generate a single CSV output file named <code>output.csv</code> (the actual file location will depend on option -o) and in the latter case the single CSV output file will be named after the <i>blend filename</i>.</p>
-e {output encoding}	<p><i>Encoding.</i> This option makes it possible to choose the character encoding of the CSV output file(s). Optional.</p> <p><i>The program defaults to UTF-8 when left to itself or when an unknown encoding is provided (in the latter case the list of encodings potentially available at runtime on the local platform is displayed as well).</i></p>
-s {output separator}	<p><i>Separator.</i> This option makes it possible to choose the field separator in the CSV output file(s). Optional.</p> <p><i>The program defaults to the semicolon when left to itself.</i></p>
-p {+ filter file}	<p><i>Positive filtering.</i> Path to an UTF-8 encoded text file (without Byte Order Mark) providing in column an explicit list of element/attribute Xpaths to extract from the XML files (while the rest will be discarded). See § 3.3.4.1 - . Optional.</p> <p><i>The program extracts all XML leaf fields when left to itself.</i> <i>Options -p and -n are exclusive.</i></p>
-n {- filter file}	<p><i>Negative filtering.</i> Path to an UTF-8 encoded text file (without Byte Order Mark) providing in column an explicit list of element/attribute Xpaths to discard from the</p>

Parameter	Description
	<p>XML files (while the rest will be extracted). See § 3.3.4.2 - . Optional.</p> <p><i>The program extracts all XML leaf fields when left to itself.</i></p> <p><i>Options <code>-n</code> and <code>-p</code> are exclusive.</i></p>
<code>-l{log4J config}</code>	<p>Logging. This option activates Log4J logging. Optional.</p> <p>Option <code>-l</code> might be immediately followed (that is: without any gap) by the path to a custom UTF-8 encoded Log4J file (without Byte Order Mark) to use instead of the default built-in Log4J configuration file <code>xml2csvlog4j.properties</code>.</p>  <p><code>xml2csvlog4j.properties</code></p> <p>Option <code>-l</code> provided alone activates Log4J logging using the built-in <code>xml2csvlog4j.properties</code> configuration file which creates a log file named <code>XML2CSV-Generic-Converter.log</code> in the same directory as the Jar file.</p>
<code>-c{limit}</code>	<p>Cutoff. This option activates cutoff limit, that is, a row count threshold for all CSV output files. When the threshold is active the program splits automatically output files if it is needed (additional CSV files are suffixed by <code>-2</code>, <code>-3</code>, ...). Optional.</p> <p>Option <code>-c</code> might be immediately followed (that is: without any gap) by an integer to use instead of the default built-in <code>1024</code> limit value.</p> <p>The actual row count threshold is <code>1024x{limit}</code> rows per CSV file (that is, with the default limit, <code>1024x1024</code> lines max for each CSV output file).</p>
<code>-w</code>	<p>Warding. Performs <i>name space</i> aware parsing. The default behavior is to confidently ignore XML elements' <i>name space</i> aliases (that is, for instance, to read <code>tagName</code> when parsing <code><lib:tagName></code>).</p> <p>When this option is set:</p> <ul style="list-style-type: none"> the actual <i>name space</i> list is recalled atop of the CSV output (before the CSV column names, displaying something like "NAMESPACES: {alias1}={namespace1} {alias2}={namespace2} ..."); each element is fully qualified by its prefix, if any (namely one of the previous <i>name space</i> aliases), and the CSV column names reflect this, keeping separating colons between an element's prefix and the element's short name (example: <code>Root.lib1:Row.lib2:Data.lib3:Tag</code>). <p>Optional.</p> <p>You won't need <i>name space</i> aware parsing unless your local XML contains <i>lookalikes</i> which must be sorted out (that is, two <i>children</i> tags of the same <i>composite</i> tag having the same name but different purposes and, as such, issued from different libraries – for instance <code><doc:Name></code> and <code><customer:Name></code>).</p> <p> Warning: <i>space full parsing wards off lookalike issues but obfuscates CSV column naming and complicates filter file's syntax as well (with the previous example: <code>Root.lib1:Row.lib2:Data.lib3:Tag</code> instead of <code>Root.Row.Data.Tag</code> both as a column name and as a valid filter XPath entry).</i></p>
<code>-u</code>	<p>Unleashed. Performs <i>root</i>-inclusive optimization. The default behavior is to exclude the <i>root</i> XML tag from the optimization process</p>

Parameter	Description
	<p>(no matter which flavor is used) for performance consideration but at the cost of a "root edge effect" (see § 3.3.3.2 -). Optional.</p> <p> CAUTION: do not unleash optimization when you handle big XML files containing root level attributes or leaf elements under the root tag at the risk of an <code>OutOfMemoryError</code>.</p> <p>Options <code>-u</code> and <code>-r</code> are exclusive.</p>

4.2 - Return codes

You won't really care about the return codes listed below unless you decide to include an `XML2CSVGenericConverter` call in a command script (for instance: a bash Unix shell script) in connection with some kind of automated process.

It is something which concerns more advanced users; if it happens that you don't know what to do with those return codes you might as well just skip this paragraph without any regret.

Return code	Description
0	<i>OK</i> . Termination without error.
1	<i>Canceled</i> , for instance because the help was displayed or because there was no XML file to process. <i>It is not an error but the program did not generate anything either.</i>
2	<i>Aborted</i> , because of some bad parameters (see § 4.1 -).
3	<i>Aborted</i> , because of filter file issues (see § 3.3.4 -).
4	<i>Aborted</i> , because the XML structure analysis failed (see § 3.3.5 -).
5	<i>Aborted</i> , because XML data extraction failed.
6	<i>Aborted</i> , because of an unexpected exception/problem.

5 - XML2CSVGenericGenerator

`XML2CSVGenericGenerator` is the core programmatic class one programmer would use in order to include a native `XML2CSVGenericConverter` call in his/her program as an internal Java routine execution.

This paragraph will be pointless for end users who are just interested in command line execution.

Programmers might find it interesting in conjunction with the Javadoc available in the `doc` directory packaged with the Jar, and the source code of the JUnit test class I wrote.

The enclosing `XML2CSVConsoleCommand` class shipped as the main class of the Jar package just:

- reads a bunch of command line options thanks to an internal `GetOpt` call and then,
- controls them, and, if it happens that the parameters are correct,
- creates an appropriate `XML2CSVGenericGenerator` class instance and, finally,
- calls its `generate` method once, with the list of XML input files to process in a row.

Appropriate logging behavior (that is: mute execution, regular logging or extra verbose/debug message display) is delegated to a `LoggingFacade` abstract class which nests a few `public static` variables in relation, making it possible to control runtime logging depth from outside the generator.

5.1 - XML2CSVGenericGenerator constructors

5.1.1 - File oriented constructors

There are two file-oriented `XML2CSVGenericGenerator` instance constructors:

```
1 // XML2CSV output file oriented constructor.
2 public XML2CSVGenericGenerator(java.io.File csvOutputFile,
3                               java.io.File csvOutputDir) throws XML2CSVException
```

This is a convenience constructor, equivalent to an `XML2CSVGenericGenerator(csvOutputFile, csvOutputDir, null, null, null, -1)` call.

```
4 // XML2CSV output file oriented constructor.
5 public XML2CSVGenericGenerator(java.io.File csvOutputFile,
6                               java.io.File csvOutputDir,
7                               java.lang.String csvFieldSeparator,
8                               java.nio.charset.Charset encoding,
9                               XML2CSVOptimization level,
10                              long cutoff) throws XML2CSVException
```

This constructor builds an `XML2CSVGenericGenerator` instance which:

- uses its `csvOutputFile` parameter as CSV output target if left non `null`;
- uses its `csvOutputDir` parameter as CSV output directory if left non `null` (letting the program name CSV output file(s) automatically after the XML input file names).

When both `csvOutputFile` and `csvOutputDir` are provided `csvOutputFile` takes precedence.

The three next parameters `csvFieldSeparator`, `encoding` and `level` pilot CSV output generation so that:

- `csvFieldSeparator` be used as CSV field separator in the output if left non `null` instead of the default character (a semicolon);
- `encoding` be used as output charset if left non `null` instead of the default UTF-8 character encoding;
- a custom optimization `level` be used instead of the default optimization level (standard - see § 3.3.3 -) if left non `null`.

The actual logging configuration used at runtime (mute execution, verbose or debug levels activated/deactivated,

built-in console logging or Log4J logging) is inherited from the current `LoggingFacade` configuration (see § 5.5 -).

The last parameter, `cutoff`, when strictly positive, activates a built-in behavior which slices automatically any output file in parts so that no output file hold more than `cutoff` lines (adding automatically `-2`, `-3`, ... suffixes to output filenames).

Just after a successful construction, a new `XML2CSVGenericGenerator` instance is ready to use (marked "*output ready*") and its `generate` method might be called immediately afterward.

5.1.2 - OutputStream oriented constructors

There are two stream-oriented `XML2CSVGenericGenerator` instance constructors:

```
11 // XML2CSV output stream oriented constructor.
12 public XML2CSVGenericGenerator(java.io.OutputStream output) throws XML2CSVException
```

This is a convenience constructor, equivalent to a `XML2CSVGenericGenerator(output, null, null, null)` call.

```
13 // XML2CSV output stream oriented constructor.
14 public XML2CSVGenericGenerator(java.io.OutputStream output,
15                               java.lang.String csvFieldSeparator,
16                               java.nio.charset.Charset encoding,
17                               XML2CSVOptimization level) throws XML2CSVException
```

This constructor builds an `XML2CSVGenericGenerator` instance which uses its `output` parameter as CSV output target (`null` value forbidden).

This direct `output` is expected opened & ready, and will remain open in all circumstances (it's the caller's responsibility to close it properly at its own level).

The other parameters have the same purpose as those provided to the previous constructor (see § 5.1.1 -).

Just like before, after a successful construction, a new `XML2CSVGenericGenerator` instance is ready to use (marked "*output ready*") and its `generate` method might be called immediately afterward.

When a single output stream is used cutoff behavior is impossible, and thresholding is automatically deactivated.

5.2 - XML2CSVGenericGenerator generation methods

`XML2CSVGenericGenerator` defines two `generate` methods:

```
18 // XML2CSV conversion.
19 public void generate(java.io.File[] xmlInputFiles,
20                     boolean withAttributes) throws XML2CSVException
```

This is a convenience method, equivalent to a `generate(xmlInputFiles, null, null, null, withAttributes)` call.

```

21 // XML2CSV conversion.
22 public void generate(java.io.File[] xmlInputFiles,
23                     java.lang.String xmlTemplateFileName,
24                     java.lang.String[] xmlExpectedElementsXPaths,
25                     java.lang.String[] xmlDiscardedElementsXPaths,
26                     boolean withAttributes) throws XML2CSVException

```

Method `generate` will raise an `XML2CSVException` if it is called while the instance is not “*output ready*”.

A newly built `XML2CSVGenericGenerator` is always marked “*output ready*”, and method `generate` will use the CSV output configuration provided to the constructor in the first place.

When method `generate` terminates the instance is marked “*idle*” and one of the three `setOutputStream`, `setOutputFile` or `setOutputDir` methods has to be called before the instance gets “*output ready*” again (see § 5.3 -).

A direct `output` (either provided to the constructor or by an explicit `setOutputStream` call) is expected opened & ready, and will remain open when `generate` terminates (responsibility left to the caller). On the contrary, regular files (either provided to the constructor or by an explicit `setOutputFile` or `setOutputDir` call) are automatically closed.

The XML input file(s) to process are provided by the `xmlInputFiles` parameter which cannot be left `null` or empty (an `XML2CSVException` would be raised).

The `xmlTemplateFileName` parameter, when not `null`, provides the name (without path) of the XML input file which should serve as XML *template file* (see § 3.3.5.1 -).

Parameters `xmlExpectedElementsXPaths` and `xmlDiscardedElementsXPaths` might be left `null`, but cannot be both non `null` (an `XML2CSVException` would be raised).

Parameter `xmlExpectedElementsXPaths` represents a *positive* Xpath *filter* list, that is, an explicit list of element Xpaths to extract from the XML input files, while the rest of XML elements will be discarded (see § 3.3.4.1 -).

Parameter `xmlDiscardedElementsXPaths` represents a *negative* Xpath *filter* list, that is, an explicit list of element Xpaths to discard from the XML input files, while the rest of XML elements will be fetched (see § 3.3.4.2 -).

The last parameter, `withAttributes`, must be set to `true` if element attributes have to be extracted alongside parsing, or set to `false` to have them discarded (see § 3.3.4.1 -).

5.3 - Other XML2CSVGenericGenerator methods

5.3.1 - Method `setOutputStream`

When called, the direct `output` parameter provided (an anonymous `OutputStream`, supposed opened & ready) becomes the new CSV output destination of the `XML2CSVGenericGenerator` instance which is marked “*output ready*” if it wasn't ready yet.

A call to this method deactivates cutoff behavior.

5.3.2 - Method `setOutputFile`

When called, the `csvOutputFile` parameter provided (a `File` instance representing a regular file) becomes the new CSV output destination of the `XML2CSVGenericGenerator` instance which is marked "*output ready*" if it wasn't ready yet.

5.3.3 - Method `setOutputDir`

When called, the `csvOutputDir` parameter provided (a `File` instance representing a regular directory) becomes the new CSV output destination of the `XML2CSVGenericGenerator` instance which is marked "*output ready*" if it wasn't ready yet.

5.3.4 - Method `setFieldSeparator`

Just like its name suggests, this method makes it possible to change or set the field separator that this `XML2CSVGenericGenerator` instance will use in CSV output.

5.3.5 - Method `setEncoding`

Just like its name suggests, this method makes it possible to change or set the character encoding that this `XML2CSVGenericGenerator` instance will use in CSV output.

5.3.6 - Method `setOptimization`

Just like its name suggests, this method makes it possible to change or set the optimization level that this `XML2CSVGenericGenerator` instance will use in CSV output (see § 3.3.3 -).

5.3.7 - Method `setCutoff`

Just like its name suggests, this method makes it possible to change or set the cutoff limit that this `XML2CSVGenericGenerator` instance will use in CSV output.

A value below zero or equals to zero deactivates cutoff behavior. When strictly positive, the value provided represents the maximum number of lines for each output CSV file.

Cutoff behavior is deactivated when a direct `OutputStream` is used (see § 5.1.2 - and § 5.3.1 -).

5.3.8 - Method `setWarding`

The parameter provided (`true` or `false`) activates or inactivates *name space* aware parsing for this

XML2CSVGenericGenerator instance.

XML2CSVGenericGenerator performs space less parsing with the default settings unless an explicit `setWarding(true)` call is performed before actual CSV generation.

5.3.9 - Method `setUnleashing`

The parameter provided (`true` or `false`) activates or inactivates *root*-inclusive optimization for this XML2CSVGenericGenerator instance.

XML2CSVGenericGenerator performs a *root*-exclusive optimization with the default settings unless an explicit `setUnleashing(true)` call is performed before actual CSV generation.

5.4 - XML2CSVGenericGenerator exceptions

XML2CSVGenericGenerator exceptions are listed below.

Exception	Description
XML2CSVCancelException	<i>Task canceled</i> , because there was no XML file to process. Strictly speaking it is not an error but because the program did not generate anything either throwing an exception ensures that the programmer will explicitly take the circumstance into account.
XML2CSVDataException	<i>Task aborted</i> , because XML data extraction failed.
XML2CSVException	Superclass of all the other XML2CSVGenericGenerator exceptions. <i>Task aborted</i> , because of an unexpected exception/problem.
XML2CSVFilterException	<i>Task aborted</i> , because of filter file issues (see § 3.3.4 -).
XML2CSVParameterException	<i>Task Aborted</i> , because of some bad parameters (see § 4.1 -).
XML2CSVStructureException	<i>Task Aborted</i> , because the XML structure analysis failed (see § 3.3.5 -).

XML2CSVGenericGenerator constructors and generation methods will throw a generic XML2CSVException in case of error that you will at your convenience refine into a more specific exception flavor instead of the default "unexpected exception/problem".

Methods `setOutputStream`, `setOutputFile` and `setOutputDir` will throw a specific XML2CSVParameterException in case of error.

The other XML2CSVGenericGenerator methods do not throw exceptions.

5.5 - LoggingFacade abstract class

This abstract class is the internal logging facade of the program in charge of progression message display and/or error message display.

The logger's static `log` methods are transparently called by XML2CSVGenericGenerator and

XML2CSVConsoleCommand each time a new message has to be displayed.

The built-in logger displays messages to the console when its `LogginFacade.log` variable is `null` (its default value), or to the underneath Log4J `org.apache.commons.logging.Log` object when set.

Apart from `LogginFacade.log`, programmers are concerned by four other public static variables which control the whole program logging capabilities, namely:

- `LogginFacade.MUTE_MODE`: mute execution. Deactivates message logging/display when set to `true`;
- `LogginFacade.DEBUG_MODE`: when set to `true`, activates the logging/display of debug messages during execution, tailored by the current debug degree. Overridden by mute mode;
- `LogginFacade.debugDegree`: when the logging/display of debug messages is activated, controls the debug degree, from the lowest to the highest. Correct values are 1 (the default), 2 or 3.
- `LogginFacade.VERBOSE_MODE`: when set to `true`, activates the logging/display of verbose progression messages. Overridden by mute mode.

6 - Command line examples

6.1 - Plain conversion of one XML file

```
java -jar XML2CSVGenericConverter_V1.0.0.jar -i c:\test\myFile.xml
```

This command line performs XML \Rightarrow CSV conversion of `myFile.xml` and generates a resulting `myFile.csv` file in directory `c:\test`. The character encoding of the CSV file is UTF-8.

The field separator is the semicolon.

A standard XML \Rightarrow CSV conversion is performed (see § 3.3.3.2 -).

Name spaces are ignored (see § 4.1 -).

Basic progression messages are displayed on the console.

Attributes, if any, are not extracted.

6.2 - Extensive #1 conversion of several XML files with attributes

```
java -jar XML2CSVGenericConverter_V1.0.0.jar -a -xXV1 -i c:\test\input -o c:\test\output
```

This command line performs XML \Rightarrow CSV conversion of all the files from directory `c:\test\input` bearing a `.xml` extension and generates as many CSV files in directory `c:\test\output`, naming the output files after the original files, with a `.csv` extension. The character encoding of the CSV files is UTF-8.

The field separator is the semicolon.

An extensive #1 XML \Rightarrow CSV conversion is performed (see § 3.3.3.3 -).

Name spaces are ignored (see § 4.1 -).

Basic progression messages are displayed on the console.

Attributes, if any, are extracted.

6.3 - Extensive #3 conversion of one XML file with log

```
java -jar XML2CSVGenericConverter_V1.0.0.jar -a -l -x -i c:\test\myFile.xml
```

This command line performs XML \Rightarrow CSV conversion of `myFile.xml` and generates a resulting `myFile.csv` file in directory `c:\test`. The character encoding of the CSV file is UTF-8.

The field separator is the semicolon.

An extensive #3 XML \Rightarrow CSV conversion is performed with the default configuration associated with option `-x` (see § 3.3.3.5 -).

Name spaces are ignored (see § 4.1 -).

A log file named `XML2CSV-Generic-Converter.log` is created in the same directory as the Jar file (configured in the built in Log4J configuration file `xml2csvlog4j.properties`).

Attributes, if any, are extracted.

6.4 - Extensive #2 conversion of one XML file with many options

```
java -jar XML2CSVGenericConverter_V1.0.0.jar -lc:\test\mylog4j.properties -xXV2 -v  
-a -w -d -i c:\test\input\myFile.xml -bmyOutputFile.csv -o c:\test\output -s,  
-eUTF-16 -u
```

This command line performs XML \Rightarrow CSV conversion of `myFile.xml` and generates a resulting `myOutputFile.csv` file in directory `c:\test\output`. The character encoding of the CSV file is UTF-16.

The field separator is the comma.

An *unleashed* extensive #2 XML \Rightarrow CSV conversion is performed (see § 3.3.3.4 -).

Name spaces are not ignored (see § 4.1 -).

A log file named and located after the custom `mylog4j.properties` configuration file is created.

The log file records verbose and debug messages.

Attributes, if any, are extracted.

6.5 - Raw conversion of several XML files in a blend output file

```
java -jar XML2CSVGenericConverter_V1.0.0.jar -lc:\test\mylog4j.properties -r -v -a  
-i c:\test\input -b -o c:\test\output
```

This command line performs XML \Rightarrow CSV conversion of all the files from directory `c:\test\input` bearing a

.xml extension and generates one unique blend file named `output.csv` in directory `c:\test\output`. The character encoding of the CSV file is UTF-8.

The field separator is the semicolon.

A raw XML \Rightarrow CSV conversion is performed (see § 3.3.3.1 -).

Name spaces are ignored (see § 4.1 -).

A log file named and located after the custom `mylog4j.properties` configuration file is created.

The log file records verbose messages.

Attributes, if any, are extracted.

7 - Real life XML examples

7.1 - All's Well That Ends Well

Ibiblio is a public digital library (URL: <http://www.ibiblio.org>).

The page <http://www.ibiblio.org/xml/examples/shakespeare/> contains Shakespeare's plays marked up in XML.

The file below is an extract of the extensive #3 XML \Rightarrow CSV conversion result of the 1st play, "*All's Well That Ends Well*" (available for download under the name `all_well.xml`).



`all_well.csv`

If you take a close look at `all_well.xml` you'll realize that a dozen of `PLAY.ACT.SCENE.SPEECH.LINE` elements contain an initial `STAGEDIR` tag before the plain `LINE` content (look for the string "*The best wishes that can be forged*" in the XML file for an example).

`PLAY.ACT.SCENE.SPEECH.LINE` elements are treated by the program as "mixtures" (see § 3.4 -) and the associated behavior (that is, here, for each `LINE`: to extract both its children elements, if any, and its ending *leaf* part) opportunely retrieves all the relevant data.

All is well that ends well... At least in Shakespeare's world. But be aware that it might not always be the case, for there is no obvious & definitive way to render "mixtures" properly in a CSV file.

Even though elements containing both text and other elements are not forbidden in XML (see http://www.w3schools.com/schema/schema_complex_mixed.asp) I think there's always a better way (that is: less loose) to provide data.

For example in `all_well.xml` an optional `STAGEDIR` attribute in each `LINE` would have been a suitable and better alternative.

7.2 - 100 MB experiment

Xmark is an XML benchmark project (URL: <http://www.xml-benchmark.org>).

An XML file generator named `xmlgen` is available on page <http://www.xml-benchmark.org/downloads.html>. On the same page a 100 MB sample XML file named `standard` is ready for download.

Although this file is quite big `XML2CSVGenericConverter` converts it without problem because the overall data *dependency* level remains low between heavily repeated elements.

However, remember that the frontier between success and failure is tight.

For example, if you modify this file and add one single small `aching="true"` attribute to the `site.regions` tag then all the `item` blocks across all the regions (`africa`, `asia`, `australia`, `europa`, `namerica`, and `samerica`) will be processed by the same buffer. I didn't count them but I know for sure an `OutOfMemoryError` is quite inevitable in such situation ⁽¹⁰⁾, for there are many many many more than one.

7.3 - The loop is closed now

The `pom.csv` file below is the *unleashed* extensive #3 XML \Rightarrow CSV conversion result of `XML2CSVGenericConverter` 's own Maven POM file.



`pom.csv`

Compactness is obviously at its peak.

Et voilà! La boucle est maintenant bouclée.

Lochrann

¹⁰ Unless you neutralize the attribute afterwards with a dedicated negative filter file excluding it from the program's scope, of course.

XML2CSVGenericConverter and this accompanying documentation come with ABSOLUTELY NO WARRANTY.
This is free software, and you are welcome to redistribute it under certain conditions.

XML2CSVGenericConverter - Copyright 2014 Laurent Popieul

lochrann@rocketmail.com

